

HR Analytics – Employee Attrition Prediction

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

import shap
from sklearn.preprocessing import StandardScaler
```

```
In [2]: df = pd.read_csv(r"C:\Users\AAFALKAZI\OneDrive\Documents\HR-Employee-Attrition.csv")
df.head()
```

Out[2]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNur
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	

5 rows × 35 columns



In [3]: `df.shape`

Out[3]: `(1470, 35)`

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Age              1470 non-null    int64  
 1   Attrition        1470 non-null    object  
 2   BusinessTravel   1470 non-null    object  
 3   DailyRate        1470 non-null    int64  
 4   Department       1470 non-null    object  
 5   DistanceFromHome 1470 non-null    int64  
 6   Education        1470 non-null    int64  
 7   EducationField   1470 non-null    object  
 8   EmployeeCount    1470 non-null    int64  
 9   EmployeeNumber   1470 non-null    int64  
 10  EnvironmentSatisfaction 1470 non-null    int64  
 11  Gender            1470 non-null    object  
 12  HourlyRate       1470 non-null    int64  
 13  JobInvolvement   1470 non-null    int64  
 14  JobLevel          1470 non-null    int64  
 15  JobRole           1470 non-null    object  
 16  JobSatisfaction  1470 non-null    int64  
 17  MaritalStatus     1470 non-null    object  
 18  MonthlyIncome     1470 non-null    int64  
 19  MonthlyRate       1470 non-null    int64  
 20  NumCompaniesWorked 1470 non-null    int64  
 21  Over18            1470 non-null    object  
 22  OverTime          1470 non-null    object  
 23  PercentSalaryHike 1470 non-null    int64  
 24  PerformanceRating 1470 non-null    int64  
 25  RelationshipSatisfaction 1470 non-null    int64  
 26  StandardHours     1470 non-null    int64  
 27  StockOptionLevel   1470 non-null    int64  
 28  TotalWorkingYears  1470 non-null    int64  
 29  TrainingTimesLastYear 1470 non-null    int64  
 30  WorkLifeBalance   1470 non-null    int64  
 31  YearsAtCompany    1470 non-null    int64  
 32  YearsInCurrentRole 1470 non-null    int64  
 33  YearsSinceLastPromotion 1470 non-null    int64  
 34  YearsWithCurrManager 1470 non-null    int64
```

```
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

In [5]: `df.describe()`

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	EnvironmentSatisfaction	HourlyRate
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.000000	1470.000000	1470.000000
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.865306	2.721769	65.89115
std	9.135373	403.509100	8.106864	1.024165	0.0	602.024335	1.093082	20.32942
min	18.000000	102.000000	1.000000	1.000000	1.0	1.000000	1.000000	30.00000
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.250000	2.000000	48.00000
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.500000	3.000000	66.00000
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.750000	4.000000	83.75000
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.000000	4.000000	100.00000

8 rows × 26 columns



TARGET VARIABLE ANALYSIS

In [6]: `df['Attrition'].value_counts(normalize=True)*100`

Out[6]: Attrition
No 83.877551
Yes 16.122449
Name: proportion, dtype: float64

Only around 16% employees have left the organization, indicating class imbalance.

DATA CLEANING

```
In [7]: df.drop(['EmployeeNumber', 'Over18', 'StandardHours'], axis=1, inplace=True)
```

FEATURE ENGINEERING

```
In [8]: df['IncomeBand'] = pd.cut(
    df['MonthlyIncome'],
    bins=[0, 3000, 7000, 20000],
    labels=['Low', 'Medium', 'High']
)
```

```
In [9]: df[['MonthlyIncome', 'IncomeBand']].head()
```

```
Out[9]:
```

	MonthlyIncome	IncomeBand
0	5993	Medium
1	5130	Medium
2	2090	Low
3	2909	Low
4	3468	Medium

```
In [10]: for col in df.select_dtypes(include='object'):
    df[col] = LabelEncoder().fit_transform(df[col])

df.head()
```

Out[10]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EnvironmentSat
0	41	1	2	1102	2	1	2	1	1	1
1	49	0	1	279	1	8	1	1	1	1
2	37	1	2	1373	1	2	2	4	1	1
3	33	0	1	1392	1	3	4	1	1	1
4	27	0	2	591	1	2	1	3	1	1


5 rows × 33 columns



EDA SECTION

In [11]:


```
sns.countplot(x='Attrition', data=df)
plt.title("Overall Attrition Distribution")
plt.show()
```



Majority employees stay, but attrition is still significant enough to impact business.

DEPARTMENT WISE ATTRITION


```
In [12]: sns.barplot(x='Department', y='Attrition', data=df)
plt.title("Department-wise Attrition")
plt.show()
```



Sales department shows higher attrition compared to R&D and HR.

SALARY VS ATTRITION


```
In [13]: sns.boxplot(x='Attrition', y='MonthlyIncome', data=df)
plt.xticks([0,1], ['No Attrition','Attrition'])
plt.title("Monthly Income vs Attrition")
plt.show()
```



Employees with lower monthly income are more likely to leave the organization.

OVERTIME IMPACT

```
In [14]: sns.countplot(x='OverTime', hue='Attrition', data=df)
plt.title("Overtime vs Attrition")
plt.show()
```



Employees working overtime have significantly higher attrition.

EDA Summary

- Attrition is highest in Sales department
- Low income employees show higher resignation rate
- Overtime and poor work-life balance are major attrition drivers
- Promotion gaps increase employee dissatisfaction

TRAIN TEST SPLIT

```
In [15]: df['IncomeBand'] = df['IncomeBand'].cat.codes
```

```
In [16]: df['IncomeBand'].dtype  
df['IncomeBand'].unique()
```

```
Out[16]: array([1, 0, 2], dtype=int8)
```

```
In [17]: X = df.drop('Attrition', axis=1)  
y = df['Attrition']  
  
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42, stratify=y  
)
```

```
In [18]: scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)  
  
lr = LogisticRegression(max_iter=1000, class_weight='balanced', solver='lbfgs')  
lr.fit(X_train_scaled, y_train)  
  
y_pred = lr.predict(X_test_scaled)  
print("Accuracy:", accuracy_score(y_test, y_pred))  
print(confusion_matrix(y_test, y_pred))  
print(classification_report(y_test, y_pred))
```

```

Accuracy: 0.7585034013605442
[[186  61]
 [ 10  37]]
      precision    recall  f1-score   support

          0       0.95     0.75     0.84     247
          1       0.38     0.79     0.51      47

   accuracy                           0.76     294
  macro avg       0.66     0.77     0.68     294
weighted avg       0.86     0.76     0.79     294

```

The numeric features were standardized using StandardScaler. Each value now represents how many standard deviations a data point is from the mean of that feature. This scaling ensures that the Logistic Regression model converges efficiently and treats all features equally in terms of magnitude.


SHAP ANALYSIS


```
In [19]: X_train_scaled_df = pd.DataFrame(
    X_train_scaled,
    columns=X.columns
)

X_test_scaled_df = pd.DataFrame(
    X_test_scaled,
    columns=X.columns
)

explainer = shap.Explainer(lr, X_train_scaled_df)
shap_values = explainer(X_test_scaled_df)

shap.summary_plot(shap_values, X_test_scaled_df)
```





SHAP analysis shows that employees working overtime with low monthly income and fewer years at the company have a significantly higher probability of attrition, while higher income, job stability, and better work-life balance reduce resignation risk.


```
In [20]: y_pred = lr.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Model Accuracy: {:.2f}%".format(accuracy*100))
```

Model Accuracy: 84.01%

```
C:\Users\AAFALKAZI\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\utils\validation.py:2684: UserWarning: X has feature names, but LogisticRegression was fitted without feature names
  warnings.warn(
```

```
In [21]: cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['No Attrition', 'Attrition'], yticklabels=['No Attrition', 'Attrition'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



The model is biased toward predicting the majority class (No Attrition). While accuracy looks okay, it fails to capture employees at risk, which is the main goal of HR attrition prediction.